

The Role of Journaling in Modern File Systems

Group Research Project

Raheim Burkett 2210198

Raul Miller 2210179

Breanna Smith 2207630

Iyana Taylor 2209566

School of Computing and Information Technology (SCIT)

Faculty of Engineering and Computing (FENC)

University of Technology, Jamaica (UTech)

CIT3002 Operating Systems (UM4)

Tutor: Mr. Philip Smith

October 25, 2024

Table of Contents

Introduction.....	3
Explanation of journaling and how it works.....	3
Comparison of journaling vs. non-journaling file systems.....	5
Real-world scenarios where journaling prevents data loss.....	7
Impact on system performance and overhead.....	8
Conclusion.....	9
Group Contributions.....	10
References.....	11

Introduction

In the current fast-moving world of digital information, data integrity and system reliability have become key concerns; hence, the role of journaling in file systems has taken great significance. This paper expounds on how journaling has become a strong mechanism in enhancing data integrity and recovery in modern file systems. In simple terms, it is a process whereby changes are pre-recorded to files by the use of journals; hence, in case systems undergo crashes or abrupt shutdowns, they have these transaction records that refer to recovery of consistency. The potential of this research does not only dwell on how journaling works, the types of journaling methods available and the benefits over non-journaling file systems, but also brings to light the pivotal role it plays in system performance and minimization of data loss. Sections of this paper address how journaling works, a comparison with non-journaling systems, applications of journaling in the real world, and its effects upon the performance and overhead of systems. Journaling is once again underscored as an invaluable contribution to the resiliency of digital infrastructure, in the context of protection means for data in contemporary file system management.

Explanation of journaling and how it works

Journaling is a method of recording information about changes that are going to be made in a file system before those changes are actually applied. Journaling enhances data integrity and keeps track of changes not yet carried out before the actual data is written to the main storage by recording them in a separate log/journal. When the initiation of a file system is being modified the system journals that operation to a special journal file that stores metadata on a disk before the actual changes occur.

Each entry of the journal stores the file address and timestamps of the particular intended changes. A transaction consists of the steps involved, each of them has a start and end marking,

therefore in case of recovery, the time and all the finished and unfinished transactions can be assigned. Every transaction made with the journals is written in one or in several blocks. Transactions are performed by users including the creation of a file, its deletion, or the modification of data.

According to the Write-ahead logging concept, the journals are performed prior to making a commitment to the file system. This method, for instance, enables the file system to complete an operation which was only partially completed, or cancel changes if they are not finished, thus preserving the integrity of the data. Once a transaction is fully logged, it is then committed by applying the changes from the journal to the actual file system. Once the task is finished, the entry in the log is denoted as done. Checkpointing entails consistently labelling journal entries as checkpointed after they have been saved to the primary storage. The system enhances performance and minimises journal usage by recording full transactions.

There are different types of Journaling, each with a different approach to logging:

- Metadata Journaling includes information about file structure, file paths, name, and directory layout and only the metadata is journaled. Since less data is written to the journal the performance is improved.
- Ordered Journaling is a hybrid approach where metadata is journaled but only commits them after the corresponding data has been written to the main storage. This type of journaling reduces the overhead of double-writing data which ensures data integrity.
- Full Data Journaling both the contents in the file and the metadata are journaled which provides the highest level of data integrity because both the data and metadata can be accurately replayed in the event of a crash. However the performance cost is high as all file data must be logged which results in double writing because it ensures that the metadata is up-to-date in the sense of recovery.

As for the advantages of using the journal, there is the downtime recovery for journals that are made to pay attention to the recent transactions , significantly reducing downtime rather than scanning and trying to detect and correct inconsistencies after a failure in the entire structure. Next, ensuring the files are consistent before a transaction and the file system can be restored to a consistent state following a failure, journaling greatly improves data integrity. Lastly, protection against data corruption by keeping a log of changes before applying them, journaling reduces the chances of file system corruption and also ensures that transactions are fully completed .

Journaling is a highly intellectual component that improves the integrity of data and reliability in file systems ensuring that transactions are safely discarded. Understanding the benefits of journaling and how it works can guide users to make informed decisions about configuring and using file systems to best meet their needs. Journaling improves the reliability and stability of file systems in various computing environments.

Comparison of journaling vs. non-journaling file systems

Journaling file systems have revolutionised data integrity and recovery processes in modern computing. Discerning the differences between journaling and non-journaling file systems becomes more important with increasing demands for reliability in today's complex systems to assure maximum efficiency and minimum data loss.

Journaling file systems recover from a system crash or power failure quite differently than non-journaling ones. In non-journaling systems, the crash initiates a file system consistency check (fsck) that looks for errors and inconsistencies across the entire system. This process can take considerable amounts of time and is often done manually when the file system is highly corrupted (Jones, 2008). In contrast, journaling file systems maintain a log - a type of journal - of recently applied

changes before actually writing those changes to the file system. This allows them to recover quickly after a crash by referencing the journal and applying unfinished changes, often restoring the system to a consistent state in minutes rather than hours (Jones, 2008). Ultimately, the key difference lies in recovery speed—while non-journaling systems require thorough checks, journaling systems can rely on pre-logged data to resume operations quickly.

When comparing the strengths and weaknesses of journaling file systems against non-journaling ones, some significant contrasts emerge. Journaling systems offer faster recovery and improved data consistency. If a crash occurs mid-operation, the journal ensures that incomplete changes are handled effectively, significantly reducing data loss. Certain journaling modes, such as ordered journaling, also ensure that data is written to its final location before recording the change in the journal, preventing inconsistent or half-written data during crashes (Prabhakaran et al., 2005).

However, journaling systems come with performance trade-offs. Writing data twice—first to the journal and then to the disk—introduces additional overhead, particularly noticeable in large sequential data writes. Moreover, the practice of grouping multiple transactions into one can lead to 'tangled synchrony,' where one synchronous process delays others, further affecting performance (Prabhakaran et al., 2005). This trade-off is a notable disadvantage when compared to non-journaling systems, which often outperform journaling systems in write-heavy workloads by skipping the journal step altogether.

In certain scenarios, non-journaling file systems excel. For instance, in workloads dominated by sequential writes, non-journaling systems like Ext4 may perform better than journaling systems, even those leveraging advanced memory technologies such as non-volatile memory (Zhang et al., 2018). Nevertheless, the overarching advantage of journaling systems remains their focus on data consistency. Even though they may sacrifice some speed in high-write environments, they ensure that file system integrity is preserved after unexpected crashes, a benefit that cannot be overlooked (Zhang et al., 2018).

Journaling and non-journaling file systems each have their strengths and weaknesses, particularly in terms of recovery processes and performance. While journaling systems prioritise data

integrity and offer faster recovery times, they often incur performance penalties in certain scenarios. Non-journaling systems can offer higher throughput in exchange for more time-consuming and less reliable recovery techniques. Understanding these trade-offs empowers one to make informed choices about the appropriate file system to deploy under various circumstances and most specifically, where data integrity is of utmost importance.

Real-world scenarios where journaling prevents data loss

In real-world scenarios, journaling technology has proven invaluable in preventing data loss across sectors by maintaining records of system changes. Journaling continuously logs data alterations, creating a log that enables systems to recover from crashes or data corruption by rolling back to a stable state. In cloud infrastructure, companies use journaling to track every data update, providing point-in-time recovery options. For example, this capability allows businesses to revert to pre-incident data versions in the event of ransomware attacks, protecting data integrity even if files are corrupted (Databarracks, 2018).

In large-scale virtualized environments, journaling also addresses the limitations of snapshots, which can negatively impact system performance when applied broadly. Unlike snapshots that capture specific moments, journaling maintains a granular history of every change, supporting continuous data replication across sites and ensuring data is recoverable despite network disruptions or hardware failures (Security Planet, 2023). The infamous “Toy Story 2” incident in 1998, in which nearly all project data was accidentally deleted, further highlights the necessity of journaling and regular backups to safeguard irreplaceable data (Databarracks, 2018).

Impact on system performance and overhead

Journaling in file systems is critical for ensuring data integrity during crashes, but it introduces performance costs, particularly in terms of write performance and disk space overhead. These performance impacts vary depending on the type of journaling used, metadata journaling or full data journaling as well as the system's workload and hardware configuration.

Journaling file systems, such as ext3, perform additional write operations because changes are first logged in a journal before being committed to the main file system. This process leads to a reduction in write performance. Metadata journaling, as implemented in systems like ext3, logs only the changes made to the file system's metadata, resulting in moderate performance overhead. Since only metadata is written twice, the performance degradation is less severe (Prabhakaran et al., 2005). In contrast, full data journaling, which logs both file data and metadata, has a more significant impact on write throughput. This method ensures more comprehensive data protection, but it can reduce write performance by up to 30% under heavy write loads (Prabhakaran et al., 2005).

In addition to affecting write performance, journaling introduces some disk space overhead. Typically, the journal consumes about 5-10% of the total file system capacity. While this is manageable in most modern systems, it can be a concern in environments with limited storage (Prabhakaran et al., 2005). However, studies have shown that optimizing journal size and placement can mitigate some of this overhead by improving the write locality, thereby minimizing the additional disk space consumption (Kim et al., 2014). Despite this disk usage, the reliability and data integrity benefits of journaling often justify the additional space requirements, especially in critical systems.

The primary advantage of journaling is its ability to significantly reduce recovery time following system crashes or power failures. Instead of performing a full file system check, which can be time-consuming on large disks, the system can simply replay the journal to restore consistency. This drastically reduces recovery time and minimizes downtime. Journaling systems like ext3, which replay

only the necessary changes recorded in the journal, can recover much faster compared to non-journaling systems, which may require full file system scans (Prabhakaran et al., 2005).

Conclusion

In summary, journaling in file systems offers a crucial layer of protection for data integrity, enabling rapid recovery and reducing the impact of unexpected system failures. By pre-logging changes and ensuring transactions are completed before they are permanently written, journaling provides an efficient mechanism to restore systems to a stable state after crashes, particularly when compared to non-journaling alternatives. The advantages of journaling—such as minimized downtime, enhanced data reliability, and significant reductions in data corruption—are particularly valuable in modern computing environments where data stability is essential. As digital reliance grows, the importance of journaling in file systems management cannot be overstated. Will the increasing complexity of our digital infrastructure call for even more advanced data integrity solutions? Embracing and advancing journaling technologies will be essential to meet the evolving demands of data resilience and reliability in the years to come.

Group Contributions

<i>Section</i>	<i>Group Member</i>
Explanation of journaling and how it works	Breanna Smith
Comparison of journaling vs. non-journaling file systems	Iyana Taylor
Real-world scenarios where journaling prevents data loss	Raheim Burkett
Impact on system performance and overhead	Raul Miller

References

Jones, T. (2008). *Anatomy of Linux journaling file Systems*. IBM Corporation.

<https://download.boulder.ibm.com/ibmdl/pub/software/dw/linux/l-journaling-fileSystems/l-journaling-fileSystems-pdf.pdf>

Prabhakaran, V., Arpaci-Dusseau, A., & Arpaci-Dusseau, R. (2005). Analysis and evolution of journaling file systems. *USENIX Annual Technical Conference*.

https://www.usenix.org/legacy/event/usenix05/tech/general/full_papers/prabhakaran/prabhakaran.pdf

Databarracks. (2018). The role of snapshots and journaling in strong data protection. Retrieved from

<https://www.databarracks.com>

Security Planet. (2023). 12 Data Loss Prevention Best Practices (+ Real Success Stories). Retrieved from <https://www.esecurityplanet.com>

Databarracks. (2018). The role of snapshots and journaling in strong data protection. Retrieved from

<https://www.databarracks.com>

Security Planet. (2023). 12 Data Loss Prevention Best Practices (+ Real Success Stories). Retrieved from <https://www.esecurityplanet.com>